

Real-Time Digital Signal Processing

Speech Enhancement Project Report

Archit Sharma: 01199766

Tomasz Bialas: 01205145

Contents

1	Introduction	3
2	Minimum Noise Buffer	3
3	Enhancements	4
3.1	Input filtering for noise estimation	4
3.1.1	Low Pass Filter, magnitude	4
3.1.2	Low Pass Filter, Power	4
3.1.3	Band Pass Filter	4
3.2	Alternative Noise Reduction Formulae	4
3.3	Oversubtraction	5
3.4	Residual Noise Reduction	5
3.5	Frame Length	6
3.6	Interval Length	6
4	Final Implementation	6
5	Testing	9
5.1	Method	9
5.2	Results	9
5.3	Performance	11
6	Conclusion	11
	Appendices	12
A	Results/Spectrograms	12
B	Code	15

1 Introduction

This project aims to implement a speech enhancer, using the DSK6713 DSP board in real-time. The presented technique will be based on the minimum noise buffering, which estimates the noise present by finding the minimum magnitudes of frequency bins over a given time period. This is then removed from the current input signal. More details of this process are discussed in Section 2. Modifications and extra enhancements made on top of this implementation are discussed and reviewed in Section 3, after which a final implementation is built and tested.

To test the speech enhancer, an audio file from SCRIBE is edited with various noises applied, and then, the speech enhancer is used to see how closely the signal is transformed into the original recording.

2 Minimum Noise Buffer

The technique used to estimate the noise of the signal is known as *spectral subtraction* [1]. It assumes a noisy signal $X(\omega)$ is equal to a clean signal $Y(\omega)$ plus some noise $N(\omega)$. Therefore, to return $Y(\omega)$, an estimate for $N(\omega)$ must be found first and then spectral subtraction $X(\omega) - N(\omega)$ performed to obtain $Y(\omega)$.

To find $N(\omega)$, this technique takes advantage of the fact that a speech signal is being processed, and thus assumes that within a time interval of approximately 10s, the speaking person will pause to take a breath (i.e. $Y(\omega) \approx 0$ when the speaker breathes, $X(\omega) \approx N(\omega)$ and if $N(\omega)$ is assumed to be fairly constant, then $N(\omega)$ can be approximated for the time interval at this point). Determining the point at which the speaker takes a breath is done by finding the magnitude of each frequency bin for each input frame and by finding the minimum magnitude for each frequency bin over the time interval.

While this method would work, it requires the DSP to store the magnitude spectrum of each frame over the 10s period. A frame is processed over a 32 ms so 312.5 magnitude spectra need to be stored as 32-bit floats, and if an FFT of size 256 is used then the memory needed is 320kB. A way to reduce this memory requirement is to split the 10s interval into four 2.5s intervals and store the smallest recorded magnitude per frequency bin over these smaller intervals. After 2.5s, the buffers rotate, with the oldest buffer then getting overwritten by the next input frame. $N(\omega)$ is then estimated by selecting the magnitude spectrum with the smallest total magnitude across all frequency bins from the four buffers. This method provides a decent estimation of $N(\omega)$ while storing only 4kB of data, a considerable saving.

The second part of the process is to subtract $N(\omega)$ from $X(\omega)$. However, the phase of the noise is unknown and instead, the magnitudes are subtracted only:

$$Y(\omega) = X(\omega) \times \left(1 - \frac{|N(\omega)|}{|X(\omega)|}\right) = X(\omega) \times g(\omega), \quad (1)$$

where $g(\omega) = \max\{\lambda, 1 - \frac{|N(\omega)|}{|X(\omega)|}\}$ to prevent negative values. It is important to note that broadband noise is random and may have peaks at random positions in the spectrum, which change location at every frame. As such, zeroing the spectrum may cause large valleys in the spectrum, causing "musical noise" with rapidly changing notes. In order to limit this effect, a small constant λ is introduced to prevent very high valleys and peaks in the spectrum. Typically, $\lambda = 0.01$ to $\lambda = 0.1$. As $N(\omega)$ is the estimated minimum noise over a period, it is likely that the noise applied to $X(\omega)$ at the current frame will be greater than the minimum being used, so $N(\omega)$ is then scaled by a constant α , initially set to 20.

The baseline algorithm was implemented and tested qualitatively in order to compare any enhancements and modifications performed. The baseline approach was found to have good a broadband noise removal, at the cost of high levels of musical noise added.

3 Enhancements

Several methods were used to improve the noise reduction of the original algorithm are described in Berouti et al., (1979) [1]. While preserving an acceptable quality of speech, only a few of those were selected for the final implementation. This was caused by the limited processing resources of the DSP and the fact that execution must be kept in real-time. Additionally, some enhancements did not improve the results and degraded the final solution.

3.1 Input filtering for noise estimation

3.1.1 Low Pass Filter, magnitude

An additional array containing the low pass filtered version of the input array was added. A pointer for the input to the noise estimation algorithm was also added, such that the estimate is performed on the low-pass version. This new array was computed based on the following formula:

$$P_t(\omega) = (1 - k) \times |X(\omega)| + k \times P_{t-1}(\omega), \quad (2)$$

where $k = \exp(-T/\tau)$. This method allowed for a noticeably improved noise suppression, and allowed to decrease the α value from 20 to 5. As a side effect, however, musical noise was added in the system, most likely resulting from the higher level difference between speech and the suppressed parts of the spectrum. This method was retained for the final implementation of the speech enhancement system.

3.1.2 Low Pass Filter, Power

A second version of the low pass filter enhancement was added, which instead of applied the formula to the magnitude spectrum $X(\omega)$, is applied to $X(\omega)^2$. The values are then converted back to magnitude values for further processing, as square root is computed.

Filtering in the power spectrum was expected to yield superior results to magnitude, as it should allow for more effective estimation of the noise and thus stronger suppression. During testing, better noise reduction was observed, however higher amounts of musical noise and heavy distortion were applied to the voice of the speaker. It was effectively "drowned out" by the musical noise. Due to those effects, this method was not selected for the final implementation.

3.1.3 Band Pass Filter

An attempt at bandpass filtering for more effective noise reduction was made by using the spectrum from approximately 300 Hz to 3.4 kHz, as typically used in telephone systems. As no improvement was achieved and the levels of musical noise resulted in unintelligible rendering of the speech, this method was not retained for the final implementation.

3.2 Alternative Noise Reduction Formulae

The initial noise removal is done by subtraction of the estimated noise magnitude through the formula $g(\omega) = \max\{\lambda, 1 - \frac{|N(\omega)|}{|X(\omega)|}\}$. In this section, several alternative formulae are implemented, and compared against the reference original implementation.

In the test code during the project development, a large `switch` statement was used to switch between different formulas, and compare their effects when integrated to the system.

$$g(\omega) = \max \left\{ \lambda \frac{|N(\omega)|}{|X(\omega)|}, 1 - \frac{|N(\omega)|}{|X(\omega)|} \right\} \quad (3)$$

$$g(\omega) = \max \left\{ \lambda \frac{|P(\omega)|}{|X(\omega)|}, 1 - \frac{|N(\omega)|}{|X(\omega)|} \right\} \quad (4)$$

$$g(\omega) = \max \left\{ \lambda \frac{|N(\omega)|}{|P(\omega)|}, 1 - \frac{|N(\omega)|}{|P(\omega)|} \right\} \quad (5)$$

$$g(\omega) = \max \left\{ \lambda, 1 - \frac{|N(\omega)|}{|P(\omega)|} \right\} \quad (6)$$

An additional switch has been implemented to test those formulae, replacing noise estimates in magnitude to estimates in the power domain, as calculated in the array mentioned in section 3.1. However, due to increased musical noise, crackle and occasional voice distortion, power domain noise estimations were not selected for the final implementation.

After testing, the equation (5) was identified as ineffective and subjectively marked as worse than the baseline in terms of noise spectrum removal. This may be explained by its use of the low-pass noise estimate as the denominator of the attenuation equation, which gives a value that is a function of those quantities, and is not directly related to the original signal and the amount of noise that requires removal. This equation did not cause musical noise or distortion but low noise rejection rendered it unfit for purpose.

The equation (6), while constructed similarly to (5), was found to have a slightly higher noise removal effect on the output. Broadband noise remained high, and additional musical noise was introduced.

Finally, the equations (3) and (4) were found to give adequate wideband noise reduction with bearable amounts of musical noise introduced. They also showed lower levels of musical noise compared to the baseline. However, those implementations introduced a decrease in quality of the speech as well as crackling into the sound.

3.3 Oversubtraction

3.4 Residual Noise Reduction

Residual noise reduction technique, as described in Boll, (1979) [2], is said to reduce musical noise. The process shown by the equation (7) states that if signal after processing $X(\omega')$ is less than some threshold η , then the output is the minimum of $X(\omega')$ across three adjacent frames - the current frame, the past frame and the next frame. This means that the output must be delayed by minimum of one frame to do this process. However, the delay of at least 32ms would not be hugely noticeable to the human ear.

$$Y(\omega)_t = \begin{cases} X(\omega')_t, & X(\omega')_t \geq \eta \\ v & \text{where } v = \min\{X(\omega')_{t-1}, X(\omega')_t, X(\omega')_{t+1}\}, \quad X(\omega')_t < \eta. \end{cases} \quad (7)$$

Because phase is being ignored, v is equal to some $X(\omega')$ which has the minimum magnitude across the three frames. Residual noise reduction works by taking advantage of the random nature of musical noise. It jumps between frequencies very quickly, so by observing adjacent frames, the musical noise will be at different frequencies. As long as the changes in the musical noise frequency are large enough to apply to different frequency bins and it does not overlap with the frequencies apparent in the clean signal, the minimum should be found and the musical noise reduced to a more consistent and less annoying hum. This decreases the previously mentioned risk of having large peaks and valleys in the spectrum, thus reducing the musical noise.

The residual noise reduction considerably reduced the amount of musical noise from the output when compared with the baseline implementation and thus will be considered for the final implementation.

3.5 Frame Length

Frame length has been set by default to 256 bins. An attempt has been made to increase the frame length to 512 as suggested to reduce musical noise, at the risk of having slurred-sounding speech. However, during testing, it has been shown that sound was heavily slurred, and that the higher size of the frame caused the CPU load to increase too much, causing skipped frames and crackling at the output. For those reasons, the frame length was left unchanged for the final implementation.

3.6 Interval Length

In order to decrease processing time and to make the system faster responding to noise, the processing interval length has been decreased from the default value of 10s to 5s or 2.5s. However, this caused increased music noise, voice distortion and crackling, and an overall heavily impacted noise reduction performance.

The processing interval was left unchanged for the final implementation.

4 Final Implementation

When testing all the enhancements, a test bench with switches that could be used in the CCS debugger was created. Only the code implementing the enhancements described in Section 3 is showcased here, the full code can be found in Appendix B.

In Figure 1, `fftframe` is the input signal after an FFT, `process_buf` points to the buffer in which processing will next occur and the other buffers store power and magnitude of the input signal before and after low pass filtering. Of note is that the `cabs` function is deconstructed here and the stages during that calculation are stored in an appropriate buffer. This is to save cycles later as some of the noise reduction formulae in Section 3.2 use these buffers to when calculating the output signal.

```

1 //magnitude filter
2 if (filtering == 1)
3 {
4     for (k=0;k<FFTLLEN;k++)
5     {
6         pow_buf[k] = fftframe[k].r*fftframe[k].r + fftframe[k].i*fftframe[k].i;
7         mag_buf[k] = sqrtf(pow_buf[k]);
8         lpmag_buf[k] = (1-filter_constant)*mag_buf[k] + filter_constant*lpmag_buf[k];
9         lppow_buf[k] = lpmag_buf[k]*lpmag_buf[k];
10    }
11    process_buf = lpmag_buf;
12 }
13
14 //power filter
15 else if (filtering == 2)
16 {
17     for (k=0;k<FFTLLEN;k++)
18     {
19         pow_buf[k] = fftframe[k].r*fftframe[k].r + fftframe[k].i*fftframe[k].i;
20         mag_buf[k] = sqrtf(pow_buf[k]);
21         lppow_buf[k] = (1-filter_constant)*pow_buf[k] + filter_constant*lppow_buf[k];
22         lpmag_buf[k] = sqrtf(lppow_buf[k]);
23    }
24    process_buf = lpmag_buf;
25 }
26
27 //no filtering *BASELINE*
28 else
29 {
30     int temp = FFTLEN;
31     for (k=0; k < temp; k++)
32     {
33         pow_buf[k] = fftframe[k].r*fftframe[k].r + fftframe[k].i*fftframe[k].i;
34         mag_buf[k] = sqrtf(pow_buf[k]);
35    }
36    process_buf = mag_buf;
37 }

```

Figure 1: Low pass filter code

To implement the minimum noise buffer rotation, the code in Figure ?? is used. The buffers rotate when `frame_count` reaches 78 as with frame processing times of 32 ms and a buffer interval time of 2.5 s, 78.125

frames are needed per buffer interval, which is then rounded down to 78. When the buffers rotate, the new buffer being written to is overwritten by the spectrum of the current frame (line 18) to prevent any information from beyond the wanted time interval of 10s impacting the speech enhancer. However, this means that the period of time from which the noise is estimated actually varies linearly between 7.5s and 10s.

```

1  if (frame_count >= 78)
2  {
3
4      /*rotates the minimum noise buffers*/
5      float *temp = m4;
6      m4 = m3;
7      m3 = m2;
8      m2 = m1;
9      m1 = temp;
10     frame_count = 0;
11
12     //shift buffer magnitude sums
13     mag4 = mag3;
14     mag3 = mag2;
15     mag2 = mag1;
16
17     //overwrites m1 with first frame in new interval
18     for (k=0;k<FFTLLEN;k++)
19     {
20         m1[k] = process_buf[k];
21     }
22 }
23 //updates m1
24 else
25 {
26     for (k=0;k<FFTLLEN;k++)
27     {
28         if (m1[k] > process_buf[k]) m1[k] = process_buf[k];
29     }
30 }
31

```

Figure 2: Buffer rotation and update code

Figure 3 shows how the minimum buffer is determined. First it calculates the magnitude sum of `m1` as it will very often change due to it updating per frame. The magnitude sums of the other buffers are fixed after the buffer is no longer being written to, and so are rotated along with the buffers themselves in Figure 2 to save clock cycles. The magnitude sums are then compared and `min_buf` then points to the minimum buffer. The power of that buffer is then calculated and stored in `noise_pow`.

```

1  //calculate m1 magnitude sum
2  mag1=0;
3  for (k=0;k<FFTLLEN;k++)
4  {
5      mag1 += m1[k];
6  }
7
8  //finds buffer with smallest magnitude sum, points min_buf to that buffer
9  if (mag1 < mag2 && mag1 < mag3 && mag1 < mag4) min_buf = m1;
10 else if (mag2 < mag3 && mag2 < mag4) min_buf = m2;
11 else if (mag3 < mag4) min_buf = m3;
12 else min_buf = m4;
13
14 //finds power buffer of that minimum noise buffer
15 for (k=0;k<FFTLLEN;k++)
16 {
17     noise_pow[k] = min_buf[k]*min_buf[k];
18 }

```

Figure 3: Minimum buffer code

The code in Figure 4 show cases how noise reduction formulae in 3.2 were switched between. Depending on `g_pow` and `g_omega`, `lambda_bot`, `lambda_top`, `signal_top` and `signal_bot` point to buffers for the appropriate formulae. In the case of λ not being multiplied by anything, `lambda_bot` and `lambda_top` point to a buffer filled with ones to keep the code readable. The reasoning behind this method becomes apparent in Figure 5.

```

1 //points lambda and signal pointers to relevent buffers depending on g_omega and g_pow
2 if (g_omega == 1)
3 {
4     lambda_top = min_buf;
5     lambda_bot = mag_buf;
6     if (g_pow)
7     {
8         signal_top = noise_pow;
9         signal_bot = pow_buf;
10    }
11    else
12    {
13        signal_top = min_buf;
14        signal_bot = mag_buf;
15    }
16 }

```

Figure 4: Example of noise reduction code

Figure 5 shows how the noise reduction is implemented. By using a generic formula with `lambda_bot`, `lambda_top`, `signal_top` and `signal_bot`, the code is very readable despite whilst being able to switch between formulae. Of note is that when `residual` is on, the processed signal is also stored in `resframe0`.

```

1 //calculates noise reduction in power domain
2 if (g_pow)
3 {
4     for (k=0;k<FFTLN;k++)
5     {
6         float temp = sqrtf(1 - alpha*(signal_top[k]/signal_bot[k]));
7         float lambda_k = lambda*(lambda_top[k]/lambda_bot[k]);
8         if (temp < lambda_k) temp = lambda_k;
9         fftframe[k].r = fftframe[k].r*temp;
10        fftframe[k].i = fftframe[k].i*temp;
11
12        //puts processed signal in resframe0 and resframe0_mag
13        if (residual)
14        {
15            resframe0[k] = fftframe[k];
16            resframe0_mag[k] = cabs(fftframe[k]);
17        }
18    }
19 }
20 }
21
22 //calculates noise reduction in magnitude *BASELINE*
23 else
24 {
25     for (k=0;k<FFTLN;k++)
26     {
27         float temp = 1 - alpha*(signal_top[k]/signal_bot[k]);
28         float lambda_k = lambda*(lambda_top[k]/lambda_bot[k]);
29         if (temp < lambda_k) temp = lambda_k;
30         fftframe[k].r = fftframe[k].r*temp;
31         fftframe[k].i = fftframe[k].i*temp;
32
33        //puts processed signal in resframe0 and resframe0_mag
34        if (residual)
35        {
36            resframe0[k] = fftframe[k];
37            resframe0_mag[k] = cabs(fftframe[k]);
38        }
39    }
40 }

```

Figure 5: Noise reduction code

Figure 6 shows how residual noise reduction is implemented. `resframe1` is the central frame, delayed by one frame with `resframe0` relatively one frame ahead and `resframe2` relatively one frame behind. The reasoning for this is described in Section 3.4.


```

1 //calculates and removes residual noise
2 if (residual)
3 {
4     complex *temp = resframe2;
5     float *temp_mag = resframe2_mag;
6     for (k=0;k<FFTLLEN;k++)
7     {
8         //checks frequency bin against threshold
9         if (resframe1_mag[k] < res_thresh)
10        {
11            //finds minimum in adjacent frames
12            if (resframe1_mag[k] < resframe0_mag[k] && resframe1_mag[k] < resframe2_mag[k]) fftframe[k] =
13            resframe1[k];
14            else if (resframe2_mag[k] < resframe0_mag[k]) fftframe[k] = resframe2[k];
15        }
16        //rotate residual buffers
17        resframe2 = resframe1;
18        resframe1 = resframe0;
19        resframe0 = temp;
20
21        resframe2_mag = resframe1_mag;
22        resframe1_mag = resframe0_mag;
23        resframe0_mag = temp_mag;
24    }

```

Figure 6: Residual noise reduction code

$$g(\omega) = \max(\lambda N(\omega), 1 - N(\omega)) \quad (8)$$

After testing using the selectable implementations, the method explained in the Berouti paper has been selected, with $\alpha = 5$, $\lambda = 0.08$, filter constant 0.7.

5 Testing

5.1 Method

Quantifying speech quality is difficult, as the perceived quality of speech can be based on many different parameters of the sound.

In order to test the speech enhancement system, two separate people would listen to the processed recording independently and then share their opinions in order to stay as objective as possible and avoid any biases. In case of conflicting opinions, judgements and justifications would be shared, and an agreement would be reached in order to match the listening test criteria as much as possible.

The design was tested against all provided sound files in order to have the widest selection of scenarios. The tested sound files were `clean`, `car1`, `factory1`, `factory2`, `lynx1`, `lynx2`, `phantom1`, `phantom2`, `phantom4`. Due to its low SNR, `phantom4` was not selected as a priority target for optimisation, as it was considered an edge case and optimisation for it caused degradation of quality in all other test cases.

During testing, a large amount of potential improvements has been rejected. Those improvements have been detailed in section 3.

Finally, an attempt to quantify the performance of the final selected design was done by recording the outputs and plotting the spectra of the different sound files.

5.2 Results

This section discusses a selected range of test results to show the system performance as well as its limits.

In the `clean` test case, the resulting spectrum was generated as expected. The spectral components of speech remained relatively unaffected, as showed by the similarity of vertical stripes in Figure 15. It can be observed that periods with no speech activity have been quieted down, even with the absence of original noise, the low background noise has been still attenuated.




Figure 7: clean file, unprocessed

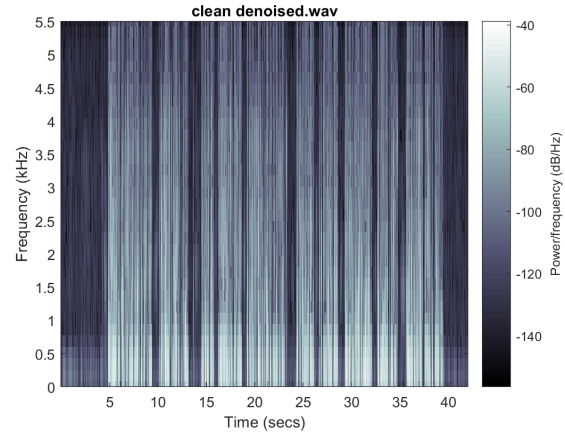


Figure 8: clean file, processed

The **factory1** test case shows the system operating in nominal conditions. The background noise has been mostly removed and periods of no voice activity can be observed on the diagram in Figure ???. Initial noise can be seen before the 5 seconds mark, after which, the noise reduction algorithm starts correctly attenuate the noise.

On the same recording, lighter sections can be seen at high frequencies at the 12, 15 and 24 seconds mark. This may be caused by sudden sounds, such as shocks and hits in the factory background noise that the system did not fully attenuate as it relies on an average noise estimate. Using statistical methods to model for such noises might be a possible solution to this issue.




Figure 9: factory1 file, unprocessed

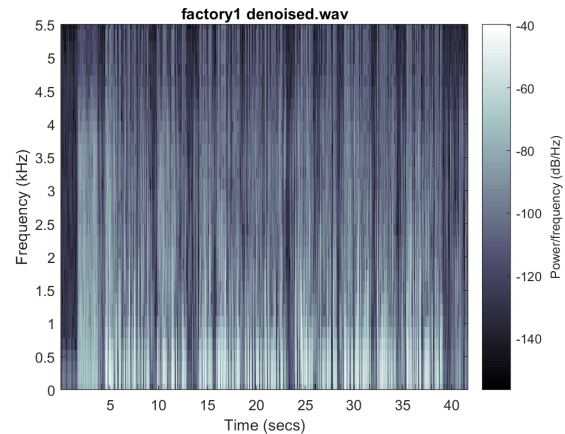


Figure 10: factory1 file, processed

Finally, in the **phantom4** test case, a degraded performance can be observed. It can be seen on the unprocessed file in Figure 30, the low SNR of the original recording. Additionally, constant noise at harmonic frequencies can be seen, which propagates to the output of the speech enhancement system, as shown in Figure ??. That periods of no speech activity still show audio activity on the output, as the system did not handle a very low SNR ratio.

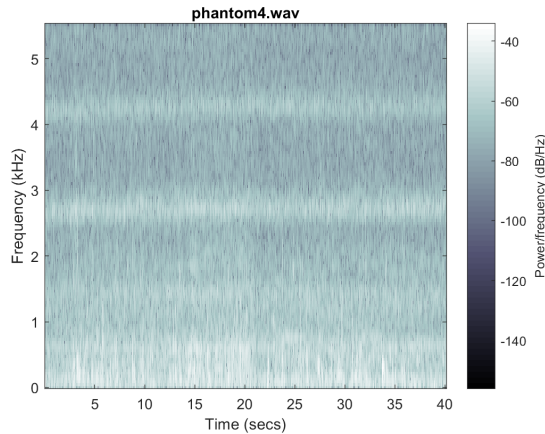


Figure 11: phantom4 file, unprocessed




Figure 12: phantom4 file, processed

However, in most cases the system outputted intelligible speech. Full test results are located in Appendix A.

5.3 Performance

Clock profiling of each enhancement was taken. When measuring the clock cycles, the numbers observed varied wildly so the results given in Figure 13 are approximate.

Enhancement	Cycles
Baseline	230000
Magnitude filtering	8000
Power filtering	47000
Residual	240000
Power Noise Reduction	70000

Figure 13: Approximate clock cycles of enhancements

From the results it becomes apparent why residual noise reduction would not work alongside the other enhancements. The number of clock cycles required to use it is considerably larger than other enhancements, and when used in conjunction with others would overrun the available CPU time.

6 Conclusion

The aim of the project was to implement a working speech enhancement system based on the spectral subtraction method and to implement, evaluate and select additional enhancements to provide optimal noise reduction performance while keeping high speech quality.

A functional system has been built, reducing background wideband noise. A compromise has been reached between speech quality, wideband noise subtraction, musical noise presence and processing time in order to meet optimum noise removal in most scenarios. Some edge cases, such as the low SNR signal in the phantom4 test case, demonstrate the limitations of the chosen design. Possible additions could include processing time optimisation, implementation of oversubtraction methods for α adjustment or the addition of the omitted residual noise reduction for performance reasons and as a result these could improve noise removal process.

Appendices

A Results/Spectrograms




Figure 14: clean file, unprocessed

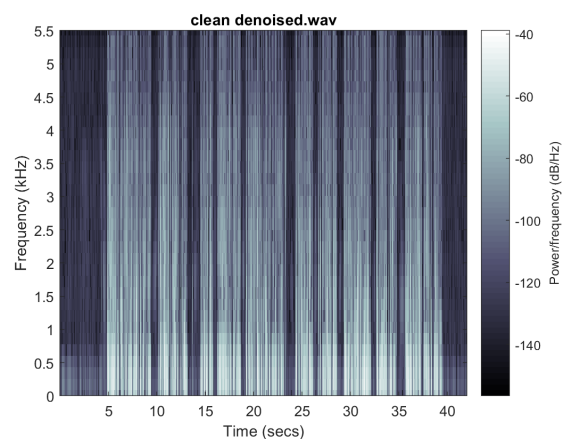


Figure 15: clean file, processed

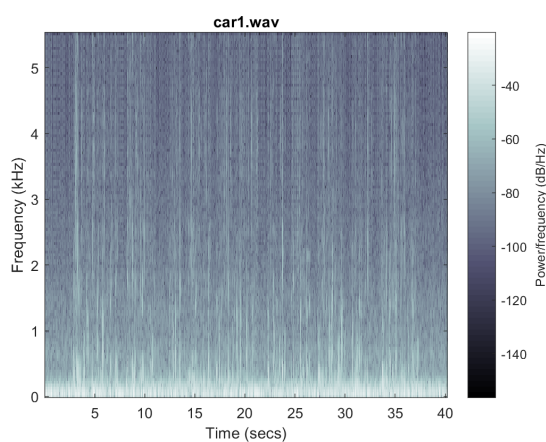


Figure 16: car1 file, unprocessed

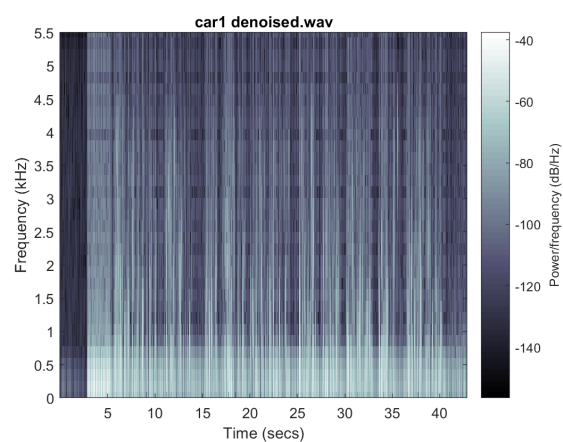


Figure 17: car1 file, processed




Figure 18: factory1 file, unprocessed




Figure 19: factory1 file, processed

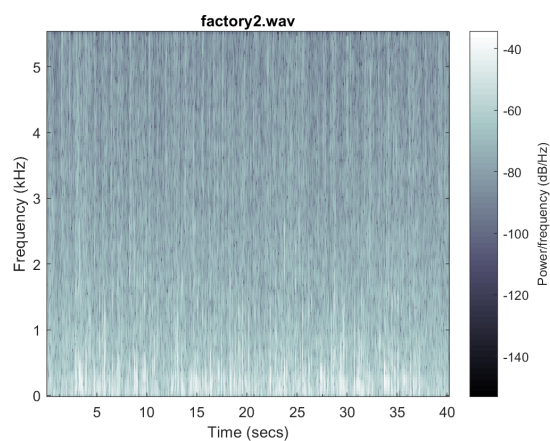


Figure 20: factory2 file, unprocessed

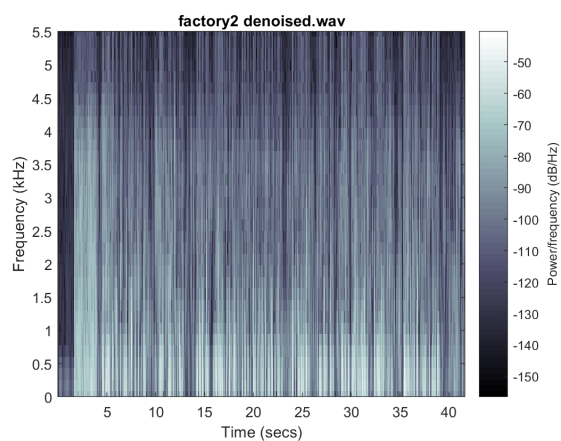


Figure 21: factory2 file, processed

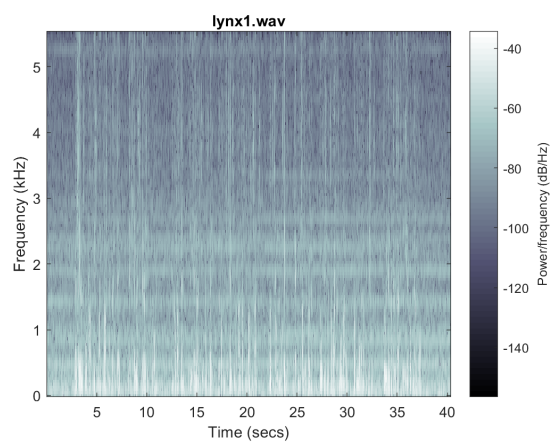


Figure 22: lynx1 file, unprocessed

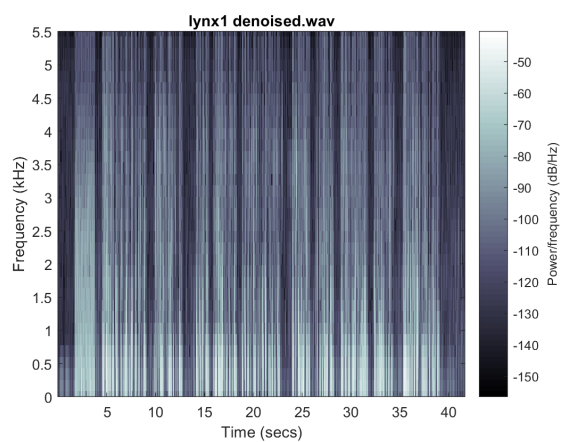


Figure 23: lynx1 file, processed




Figure 24: lynx2 file, unprocessed




Figure 25: lynx2 file, processed




Figure 26: phantom1 file, unprocessed

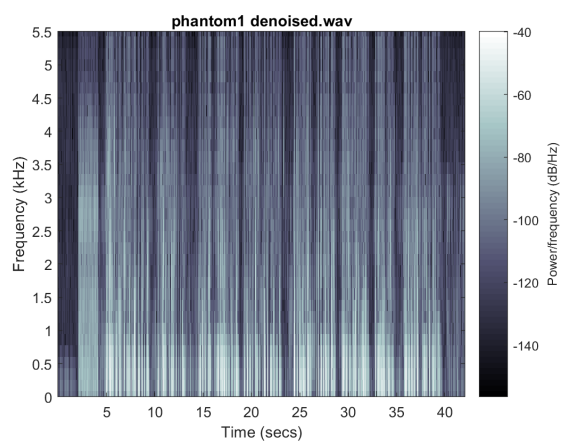


Figure 27: phantom1 file, processed




Figure 28: phantom2 file, unprocessed

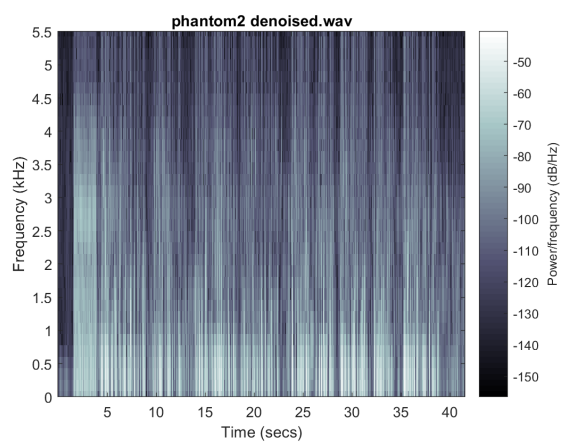


Figure 29: phantom2 file, processed




Figure 30: phantom4 file, unprocessed

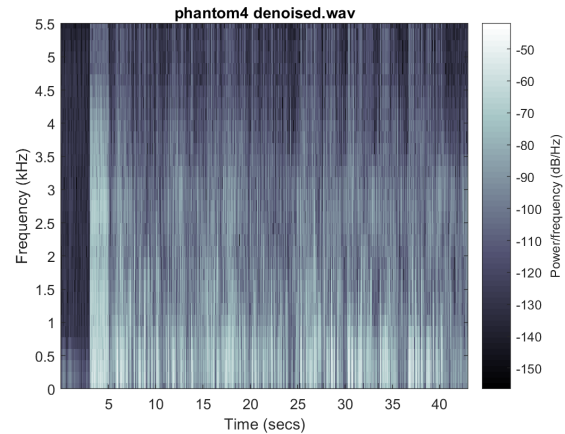


Figure 31: phantom4 file, processed

B Code

```

1 enhance_a_new_hope.c
2
3 Type
4 Text
5 Size
6 16 KB (16,678 bytes)
7 Storage used
8 0 bytes Owned by undefined
9 Location
10 project
11 Owner
12 Tomasz Bia as
13 Modified
14 Mar 20, 2019 by Tomasz Bia as
15 Opened
16 6:57 PM by me
17 Created
18 Mar 20, 2019
19 Add a description
20 Viewers can download
21
22 /*****
23  DEPARTMENT OF ELECTRICAL AND ELECTRONIC ENGINEERING
24  IMPERIAL COLLEGE LONDON
25
26  EE 3.19: Real Time Digital Signal Processing
27  Dr Paul Mitcheson and Daniel Harvey
28
29  PROJECT: Frame Processing
30
31  ***** ENHANCE. C *****
32  Shell for speech enhancement
33
34  Demonstrates overlap-add frame processing (interrupt driven) on the DSK.
35
36  By Danny Harvey: 21 July 2006
37  Updated for use on CCS v4 Sept 2010
38  *****/
39
40 /* You should modify the code so that a speech enhancement project is built
41  * on top of this template.
42  */
43
44 /***** Pre-processor statements *****/
45 // library required when using calloc
46 #include <stdlib.h>
47 // Included so program can make use of DSP/BIOS configuration tool.
48 #include "dsp_bios_cfg.h"
49
50 /* The file dsk6713.h must be included in every program that uses the BSL. This
51  example also includes dsk6713-aic23.h because it uses the
52  AIC23 codec module (audio interface). */
53 #include "dsk6713.h"
54 #include "dsk6713-aic23.h"
55
56 // math library (trig functions)
57 #include <math.h>
58
59 /* Some functions to help with Complex algebra and FFT. */
60 #include "cmplx.h"
61 #include "fft_functions.h"
62
63 // Some functions to help with writing/reading the audio ports when using interrupts.
64 #include <helper_functions_ISR.h>
65
66 #define WINCONST 0.85185 /* 0.46/0.54 for Hamming window */
67 #define FSAMP 8000.0 /* sample frequency, ensure this matches Config for AIC */
68 #define FFTLEN 256 /* fft length = frame length 256/8000 = 32 ms */
69 #define NFREQ (1+FFTLEN/2) /* number of frequency bins from a real FFT */
70 #define OVERSAMP 4 /* oversampling ratio (2 or 4) */

```

```

70 #define FRAMEINC (FFTLEN/OVERSAMP) /* Frame increment */
71 #define CIRCBUF (FFTLEN+FRAMEINC) /* length of I/O buffers */
72
73 #define OUTGAIN 64000.0 /* Output gain for DAC */
74 #define INGAIN (1.0/16000.0) /* Input gain for ADC */
75 // PI defined here for use in your code
76 #define PI 3.141592653589793
77 #define TFRAME FRAMEINC/FSAMP /* time between calculation of each frame */
78
79 /***** Global declarations *****/
80
81 /* Audio port configuration settings: these values set registers in the AIC23 audio
82 interface to configure it. See TI doc SLWS106D 3-3 to 3-10 for more info. */
83 DSK6713_AIC23_Config Config = { \
84     /* *****/
85     /* REGISTER FUNCTION SETTINGS */
86     /* *****/
87     0x0017, /* 0 LEFTINVOL Left line input channel volume 0dB */
88     0x0017, /* 1 RIGHTINVOL Right line input channel volume 0dB */
89     0x01f9, /* 2 LEFTHPVOL Left channel headphone volume 0dB */
90     0x01f9, /* 3 RIGHTHPVOL Right channel headphone volume 0dB */
91     0x0011, /* 4 ANAPATH Analog audio path control DAC on, Mic boost 20dB */
92     0x0000, /* 5 DIGPATH Digital audio path control All Filters off */
93     0x0000, /* 6 DPOWERDOWN Power down control All Hardware on */
94     0x0043, /* 7 DIGIF Digital audio interface format 16 bit */
95     0x008d, /* 8 SAMPLERATE Sample rate control 8 KHZ-ensure matches FSAMP */
96     0x0001, /* 9 DIGACT Digital interface activation On */
97     /* *****/
98 };
99
100 // Codec handle:- a variable used to identify audio interface
101 DSK6713_AIC23_CodecHandle H_Codec;
102
103 float *inbuffer, *outbuffer; /* Input/output circular buffers */
104 complex *fftframe; /* processing frame */
105 float *inwin, *outwin; /* Input and output windows */
106 float ingain, outgain; /* ADC and DAC gains */
107 float cpufract; /* Fraction of CPU time used */
108 volatile int io_ptr=0; /* Input/output pointer for circular buffers */
109 volatile int frame_ptr=0; /* Frame pointer */
110
111 float *m1, *m2, *m3, *m4; /* Minimum noise intervals */
112 float mag1 = 0; /* Magnitude sum of m1 */
113 float mag2 = 0; /* Magnitude sum of m2 */
114 float mag3 = 0; /* Magnitude sum of m3 */
115 float mag4 = 0; /* Magnitude sum of m4 */
116
117 float *noise_pow; /* Noise power for selected minimum buffer */
118 float *lpmag_buf; /* Magnitude of input after low pass filter buffer */
119 float *lppow_buf; /* Power of input after low pass filter buffer */
120 float *mag_buf; /* Magnitude of input buffer */
121 float *pow_buf; /* Power of input buffer */
122
123 complex *resframe0; /* Stores processed signal one frame ahead of resframe1 */
124 complex *resframe1; /* Stores processed signal to be output at a one frame delay from input */
125 complex *resframe2; /* Stores processed signal one frame behind of resframe1 */
126
127 float *resframe0_mag; /* magnitude buffer of resframe0 */
128 float *resframe1_mag; /* magnitude buffer of resframe1 */
129 float *resframe2_mag; /* magnitude buffer of resframe2 */
130
131 float *ones; /* buffer of ones, used to switch easily between noise reduction formulae */
132
133 float lambda = 0.08;
134 float alpha = 5.0; /* scales noise reduction */
135
136 float time_constant = 0.02; /* time constant for low pass filter */
137
138 float res_thresh = 0.8; /* residual noise threshold */
139
140 int frame_count = 0; /* counter to see when an interval is over */
141 int frames_per_interval; /* number of frames when interval is over */
142
143 float time_frame = 10.0; /* time interval over which minimum noise is found */
144 float filter_constant; /* calculated within main, declared here if desired to change manually */
145
146 /***** switches *****/
147
148 int passthrough = 0; /* allows signal through without processing */
149 int filtering = 1; /* switches filtering; 0 = no filtering, 1 = magnitude filtering, 2 = power
150 filtering */
151 int g_omega = 5; /* switches between noise reduction formulae */
152 int g_pow = 0; /* turns on noise reduction in power domain */
153 int residual = 0; /* turns on residual noise reduction */
154
155 /***** Function prototypes *****/
156 void init_hardware(void); /* Initialize codec */
157 void init_HWI(void); /* Initialize hardware interrupts */
158 void ISR_AIC(void); /* Interrupt service routine for codec */
159 void process_frame(void); /* Frame processing routine */
160
161 /***** Main routine *****/
162 void main()
163 {
164     int k; /* used in various for loops */
165
166     filter_constant = 0.7; /* expf(-TFRAME/time_constant); */
167
168     /* Initialize and zero fill arrays */
169
170     inbuffer = (float *) calloc(CIRCBUF, sizeof(float)); /* Input array */
171     outbuffer = (float *) calloc(CIRCBUF, sizeof(float)); /* Output array */
172     fftframe = (complex *) calloc(FFTLEN, sizeof(complex));
173     inwin = (float *) calloc(FFTLEN, sizeof(float)); /* Input window */
174     outwin = (float *) calloc(FFTLEN, sizeof(float)); /* Output window */
175
176     /* minimum noise buffers */
177
178

```



```

180     m1      = (float *) calloc(FFTLLEN, sizeof(float)); /* noise buffer*/
181     m2      = (float *) calloc(FFTLLEN, sizeof(float)); /* noise buffer*/
182     m3      = (float *) calloc(FFTLLEN, sizeof(float)); /* noise buffer*/
183     m4      = (float *) calloc(FFTLLEN, sizeof(float)); /* noise buffer*/
184
185     /*magnitude and power buffers of noise and signal at various stages*/
186
187     lpmag_buf = (float *) calloc(FFTLLEN, sizeof(float));
188     lppow_buf = (float *) calloc(FFTLLEN, sizeof(float));
189     mag_buf   = (float *) calloc(FFTLLEN, sizeof(float)); /* noise buffer*/
190     pow_buf   = (float *) calloc(FFTLLEN, sizeof(float)); /* noise buffer*/
191     noise_pow = (float *) calloc(FFTLLEN, sizeof(float));
192
193     /*residual noise reduction buffers*/
194
195     resframe0 = (complex *) calloc(FFTLLEN, sizeof(complex));
196     resframe1 = (complex *) calloc(FFTLLEN, sizeof(complex));
197     resframe2 = (complex *) calloc(FFTLLEN, sizeof(complex));
198
199     resframe0_mag = (float *) calloc(FFTLLEN, sizeof(float));
200     resframe1_mag = (float *) calloc(FFTLLEN, sizeof(float));
201     resframe2_mag = (float *) calloc(FFTLLEN, sizeof(float));
202
203     /*array of ones used for g-omega and g-pow switching to be easier*/
204
205     ones = (float *) calloc(FFTLLEN, sizeof(float));
206
207     /* initialize board and the audio port */
208     init_hardware();
209
210     /* initialize hardware interrupts */
211     init_HWI();
212
213     /* initialize algorithm constants */
214
215     for (k=0;k<FFTLLEN;k++)
216     {
217         inwin[k] = sqrt((1.0-WINCONST*cos(PI*(2*k+1)/FFTLLEN))/OVERSAMP);
218         outwin[k] = inwin[k];
219         ones[k] = 1; //fill ones with ones
220     }
221     ingain=INGAIN;
222     outgain=OUTGAIN;
223
224     /* main loop, wait for interrupt */
225     while(1) process_frame();
226 }
227
228 /***** init_hardware() *****/
229 void init_hardware()
230 {
231     // Initialize the board support library, must be called first
232     DSK6713_init();
233
234     // Start the AIC23 codec using the settings defined above in config
235     H_Codec = DSK6713_AIC23_openCodec(0, &Config);
236
237     /* Function below sets the number of bits in word used by MSBSP (serial port) for
238     receives from AIC23 (audio port). We are using a 32 bit packet containing two
239     16 bit numbers hence 32BIT is set for receive */
240     MCBSP_FSETS(RCR1, RWDLEN1, 32BIT);
241
242     /* Configures interrupt to activate on each consecutive available 32 bits
243     from Audio port hence an interrupt is generated for each L & R sample pair */
244     MCBSP_FSETS(SPCR1, RINTM, FRM);
245
246     /* These commands do the same thing as above but applied to data transfers to the
247     audio port */
248     MCBSP_FSETS(XCR1, XWDLEN1, 32BIT);
249     MCBSP_FSETS(SPCR1, XINTM, FRM);
250
251 }
252
253 /***** init_HWI() *****/
254 void init_HWI(void)
255 {
256     IRQ_globalDisable(); // Globally disables interrupts
257     IRQ_nmiEnable(); // Enables the NMI interrupt (used by the debugger)
258     IRQ_map(IRQ_EVT_RINT1,4); // Maps an event to a physical interrupt
259     IRQ_enable(IRQ_EVT_RINT1); // Enables the event
260     IRQ_globalEnable(); // Globally enables interrupts
261 }
262
263
264
265 /***** process_frame() *****/
266 void process_frame(void)
267 {
268     int k, m; //iteration variables
269     int io_ptr0;
270     float *min_buf; //points to minimum noise buffer
271     float *process_buf; //points to buffer to be processed
272
273     float *lambda_top; //lambda multiplication numerator for g-omega and g-pow switching
274     float *lambda_bot; //lambda multiplication denominator for g-omega and g-pow switching
275     float *signal_top; //signal multiplication numerator for g-omega and g-pow switching
276     float *signal_bot; //signal multiplication denominator for g-omega and g-pow switching
277
278     /* work out fraction of available CPU time used by algorithm */
279     cpufrac = ((float) (io_ptr & (FRAMEINC - 1)))/FRAMEINC;
280
281     /* wait until io_ptr is at the start of the current frame */
282     while((io_ptr/FRAMEINC) != frame_ptr);
283
284     /* then increment the framecount (wrapping if required) */
285     if (++frame_ptr >= (CIRCBUF/FRAMEINC)) frame_ptr=0;
286
287     /* save a pointer to the position in the I/O buffers (inbuffer/outbuffer) where the
288     data should be read (inbuffer) and saved (outbuffer) for the purpose of processing */
289     io_ptr0=frame_ptr * FRAMEINC;
290

```

```

291 | /* copy input data from inbuffer into inframe (starting from the pointer position) */
292 |
293 | m=io_ptr0;
294 | for (k=0;k<FFTLLEN;k++)
295 | {
296 |     fftframe[k].r = inbuffer[m] * inwin[k];
297 |     fftframe[k].i = 0;
298 |     if (++m >= CIRCBUF) m=0; /* wrap if required */
299 | }
300 |
301 | /***** DO PROCESSING OF FRAME HERE *****/
302 | frames_per_interval = (int)(time_frame/TFRAME);
303 | fft(FFTLLEN, fftframe);
304 |
305 | //allows signal to pass through unprocessed
306 | if(!passthrough)
307 | {
308 |
309 |     frame_count++;
310 |
311 |     //magnitude filter
312 |     if (filtering == 1)
313 |     {
314 |         for (k=0;k<FFTLLEN;k++)
315 |         {
316 |             pow_buf[k] = fftframe[k].r*fftframe[k].r + fftframe[k].i*fftframe[k].i;
317 |             mag_buf[k] = sqrtf(pow_buf[k]);
318 |             lpmag_buf[k] = (1-filter_constant)*mag_buf[k] + filter_constant*lpmag_buf[k];
319 |             lppow_buf[k] = lpmag_buf[k]*lpmag_buf[k];
320 |         }
321 |         process_buf = lpmag_buf;
322 |     }
323 |
324 |     //power filter
325 |     else if (filtering == 2)
326 |     {
327 |         for (k=0;k<FFTLLEN;k++)
328 |         {
329 |             pow_buf[k] = fftframe[k].r*fftframe[k].r + fftframe[k].i*fftframe[k].i;
330 |             mag_buf[k] = sqrtf(pow_buf[k]);
331 |             lppow_buf[k] = (1-filter_constant)*pow_buf[k] + filter_constant*lppow_buf[k];
332 |             lpmag_buf[k] = sqrtf(lppow_buf[k]);
333 |         }
334 |         process_buf = lpmag_buf;
335 |     }
336 |
337 |     //no filtering *BASELINE*
338 |     else
339 |     {
340 |         int temp = FFTLEN;
341 |         for (k=0; k < temp; k++)
342 |         {
343 |             pow_buf[k] = fftframe[k].r*fftframe[k].r + fftframe[k].i*fftframe[k].i;
344 |             mag_buf[k] = sqrtf(pow_buf[k]);
345 |         }
346 |         process_buf = mag_buf;
347 |     }
348 |
349 |     if (frame_count >= 78)
350 |     {
351 |
352 |         /*rotates the minimum noise buffers*/
353 |         float *temp = m4;
354 |         m4 = m3;
355 |         m3 = m2;
356 |         m2 = m1;
357 |         m1 = temp;
358 |         frame_count = 0;
359 |
360 |         //shift buffer magnitude sums
361 |         mag4 = mag3;
362 |         mag3 = mag2;
363 |         mag2 = mag1;
364 |
365 |         //overwrites m1 with first frame in new interval
366 |         for (k=0;k<FFTLLEN;k++)
367 |         {
368 |             m1[k] = process_buf[k];
369 |         }
370 |
371 |         //updates m1
372 |         else
373 |         {
374 |             for (k=0;k<FFTLLEN;k++)
375 |             {
376 |                 if (m1[k] > process_buf[k]) m1[k] = process_buf[k];
377 |             }
378 |
379 |             //calculate m1 magnitude sum
380 |             mag1=0;
381 |             for (k=0;k<FFTLLEN;k++)
382 |             {
383 |                 mag1 += m1[k];
384 |             }
385 |
386 |             //finds buffer with smallest magnitude sum, points min_buf to that buffer
387 |             if (mag1 < mag2 && mag1 < mag3 && mag1 < mag4) min_buf = m1;
388 |             else if (mag2 < mag3 && mag2 < mag4) min_buf = m2;
389 |             else if (mag3 < mag4) min_buf = m3;
390 |             else min_buf = m4;
391 |
392 |             //finds power buffer of that minimum noise buffer
393 |             for (k=0;k<FFTLLEN;k++)
394 |             {
395 |                 noise_pow[k] = min_buf[k]*min_buf[k];
396 |             }
397 |
398 |             //points lambda and signal pointers to relevant buffers depending on g_omega and g_pow
399 |             if (g_omega == 1)
400 |             {
401 |                 lambda_top = min_buf;

```

```

402     lambda_bot = mag_buf;
403     if (g_pow)
404     {
405         signal_top = noise_pow;
406         signal_bot = pow_buf;
407     }
408     else
409     {
410         signal_top = min_buf;
411         signal_bot = mag_buf;
412     }
413 }
414
415 else if (g_omega == 2)
416 {
417     lambda_top = lpmag_buf;
418     lambda_bot = mag_buf;
419     if (g_pow)
420     {
421         signal_top = noise_pow;
422         signal_bot = pow_buf;
423     }
424     else
425     {
426         signal_top = min_buf;
427         signal_bot = mag_buf;
428     }
429 }
430
431 else if (g_omega == 3)
432 {
433     lambda_top = min_buf;
434     lambda_bot = lpmag_buf;
435     if (g_pow)
436     {
437         signal_top = noise_pow;
438         signal_bot = lppow_buf;
439     }
440     else
441     {
442         signal_top = min_buf;
443         signal_bot = lpmag_buf;
444     }
445 }
446
447 else if (g_omega == 4)
448 {
449     lambda_top = ones;
450     lambda_bot = ones;
451     if (g_pow)
452     {
453         signal_top = noise_pow;
454         signal_bot = lppow_buf;
455     }
456     else
457     {
458         signal_top = min_buf;
459         signal_bot = lpmag_buf;
460     }
461 }
462
463 else if (g_omega == 5)
464 {
465     lambda_top = min_buf;
466     lambda_bot = ones;
467     if (g_pow)
468     {
469         signal_top = noise_pow;
470         signal_bot = pow_buf;
471     }
472     else
473     {
474         signal_top = min_buf;
475         signal_bot = mag_buf;
476     }
477 }
478 //default baseline
479 else
480 {
481     lambda_top = ones;
482     lambda_bot = ones;
483     if (g_pow)
484     {
485         signal_top = noise_pow;
486         signal_bot = pow_buf;
487     }
488     else
489     {
490         signal_top = min_buf;
491         signal_bot = mag_buf;
492     }
493 }
494
495 //calculates noise reduction in power domain
496 if (g_pow)
497 {
498     for (k=0;k<FFTLLEN;k++)
499     {
500         float temp = sqrtf(1 - alpha*(signal_top[k]/signal_bot[k]));
501         float lambda_k = lambda*(lambda_top[k]/lambda_bot[k]);
502         if (temp < lambda_k) temp = lambda_k;
503         fftframe[k].r = fftframe[k].r*temp;
504         fftframe[k].i = fftframe[k].i*temp;
505
506         //puts processed signal in resframe0 and resframe0_mag
507         if (residual)
508         {
509             resframe0[k] = fftframe[k];
510             resframe0_mag[k] = cabs(fftframe[k]);
511         }
512     }

```

```

513     }
514 }
515
516 //calculates noise reduction in magnitude *BASELINE*
517 else
518 {
519     for (k=0;k<FFTLLEN;k++)
520     {
521         float temp = 1 - alpha*(signal_top[k]/signal_bot[k]);
522         float lambda_k = lambda*(lambda_top[k]/lambda_bot[k]);
523         if (temp < lambda_k) temp = lambda_k;
524         fftframe[k].r = fftframe[k].r*temp;
525         fftframe[k].i = fftframe[k].i*temp;
526
527         //puts processed signal in resframe0 and resframe0_mag
528         if (residual)
529         {
530             resframe0[k] = fftframe[k];
531             resframe0_mag[k] = cabs(fftframe[k]);
532         }
533     }
534 }
535
536 //calculates and removes residual noise
537 if (residual)
538 {
539     complex *temp = resframe2;
540     float *temp_mag = resframe2_mag;
541     for (k=0;k<FFTLLEN;k++)
542     {
543         //checks frequency bin against threshold
544         if (resframe1_mag[k] < res_thresh)
545         {
546             //finds minimum in adjacent frames
547             if (resframe1_mag[k] < resframe0_mag[k] && resframe1_mag[k] < resframe2_mag[k]) fftframe[k] =
resframe1[k];
548             else if (resframe2_mag[k] < resframe0_mag[k]) fftframe[k] = resframe2[k];
549         }
550     }
551     //rotate residual buffers
552     resframe2 = resframe1;
553     resframe1 = resframe0;
554     resframe0 = temp;
555
556     resframe2_mag = resframe1_mag;
557     resframe1_mag = resframe0_mag;
558     resframe0_mag = temp_mag;
559 }
560
561 }
562 ifft(FFTLLEN, fftframe);
563
564 /*****
565  * multiply outframe by output window and overlap-add into output buffer */
566
567 m=io_ptr0;
568
569 for (k=0;k<(FFTLLEN-FRAMEINC);k++)
570 {
571     /* this loop adds into outbuffer */
572     outbuffer[m] = outbuffer[m]+fftframe[k].r*outwin[k];
573     if (++m >= CIRCBUF) m=0; /* wrap if required */
574 }
575 for (;k<FFTLLEN;k++)
576 {
577     outbuffer[m] = fftframe[k].r*outwin[k]; /* this loop over-writes outbuffer */
578     m++;
579 }
580 }
581 /***** INTERRUPT SERVICE ROUTINE *****/
582
583 // Map this to the appropriate interrupt in the CDB file
584
585 void ISR_AIC(void)
586 {
587     short sample;
588     /* Read and write the ADC and DAC using inbuffer and outbuffer */
589
590     sample = mono_read_16Bit();
591     inbuffer[io_ptr] = ((float)sample)*ingain;
592     /* write new output data */
593     mono_write_16Bit((int)(outbuffer[io_ptr]*outgain));
594
595     /* update io_ptr and check for buffer wraparound */
596     if (++io_ptr >= CIRCBUF) io_ptr=0;
597 }
598
599 /*****
600  *
601  */

```

References

- [1] M Berouti, Richard Schwartz and J Makhoul.
 “Enhancement of speech corrupted by acoustic noise”.
 In: *[No source information available]* 4 (May 1979), pp. 208–211.
 DOI: 10.1109/ICASSP.1979.1170788.

- [2] Steven F. Boll. “Suppression of Acoustic Noise in Speech Using Spectral Subtraction”.
In: *Acoustics, Speech and Signal Processing, IEEE Transactions on* 27 (May 1979), pp. 113–120.
DOI: 10.1109/TASSP.1979.1163209.